

# Guided Tour of Machine Learning in Finance


## **Week 2 - Lesson 1 - part 5: Gradient Descent Optimization**

Igor Halperin

NYU Tandon School of Engineering, 2017


# Linear Regression as a one-neuron calculation

- **Inputs:** real-valued numbers
- **Output:** a real-valued number
- Can be thought of as a node computing a **linear function** of inputs
- Weights can be tuned to fit the data



# Linear Regression as a one-neuron calculation

- **Inputs:** real-valued numbers
- **Output:** a real-valued number
- Can be thought of as a node computing a **linear function** of inputs
- Weights can be tuned to fit the data



- For Linear Regression, there is an **analytical (one-step) solution** via the normal equation

# Analytical learning in Linear Regression

**Performance measure P** : mean square error (**MSE**)  $\text{MSE}_{test}$  :

$$\text{MSE}_{test} = \frac{1}{N_{test}} \sum_{n=1}^{N_{test}} (\hat{y}_i^{test} - y_i^{test})^2 = \frac{1}{N_{test}} \left\| \hat{\mathbf{Y}}^{test} - \mathbf{Y}^{test} \right\|_2^2 = \frac{1}{N_{test}} \left\| \mathbf{X}^{test} \mathbf{W} - \mathbf{Y}^{test} \right\|_2^2$$

Parameters  $\mathbf{W} \in \mathbb{R}^D$  are found by minimizing  $\text{MSE}_{train}$

$$\text{MSE}_{train} = \frac{1}{N_{train}} \sum_{n=1}^{N_{train}} (\hat{y}_i^{train} - y_i^{train})^2 = \frac{1}{N_{train}} \left\| \hat{\mathbf{Y}}^{train} - \mathbf{Y}^{train} \right\|_2^2 = \frac{1}{N_{train}} \left\| \mathbf{X}^{train} \mathbf{W} - \mathbf{Y}^{train} \right\|_2^2$$

For optimal weights  $\mathbf{W}$ , the gradient of  $\text{MSE}_{train}$  should be 0:

$$\nabla_{\mathbf{W}} \text{MSE}_{train} = \nabla_{\mathbf{W}} \frac{1}{N_{train}} \left\| \hat{\mathbf{Y}}^{train} - \mathbf{Y}^{train} \right\|_2^2 = \frac{1}{N_{train}} \nabla_{\mathbf{W}} \left\| \mathbf{X}^{train} \mathbf{W} - \mathbf{Y}^{train} \right\|_2^2 = 0$$

$\Rightarrow$

$$\mathbf{W} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Normal  
equation

# Analytical learning in Linear Regression

**Performance measure P** : mean square error (**MSE**)  $\text{MSE}_{test}$  :

$$\text{MSE}_{test} = \frac{1}{N_{test}} \sum_{n=1}^{N_{test}} (\hat{y}_i^{test} - y_i^{test})^2 = \frac{1}{N_{test}} \left\| \hat{\mathbf{Y}}^{test} - \mathbf{Y}^{test} \right\|_2^2 = \frac{1}{N_{test}} \left\| \mathbf{X}^{test} \mathbf{W} - \mathbf{Y}^{test} \right\|_2^2$$

Parameters  $\mathbf{W} \in \mathbb{R}^D$  are found by minimizing  $\text{MSE}_{train}$

$$\text{MSE}_{train} = \frac{1}{N_{train}} \sum_{n=1}^{N_{train}} (\hat{y}_i^{train} - y_i^{train})^2 = \frac{1}{N_{train}} \left\| \hat{\mathbf{Y}}^{train} - \mathbf{Y}^{train} \right\|_2^2 = \frac{1}{N_{train}} \left\| \mathbf{X}^{train} \mathbf{W} - \mathbf{Y}^{train} \right\|_2^2$$

For optimal weights  $\mathbf{W}$ , the gradient of  $\text{MSE}_{train}$  should be 0:

$$\nabla_{\mathbf{W}} \text{MSE}_{train} = \nabla_{\mathbf{W}} \frac{1}{N_{train}} \left\| \hat{\mathbf{Y}}^{train} - \mathbf{Y}^{train} \right\|_2^2 = \frac{1}{N_{train}} \nabla_{\mathbf{W}} \left\| \mathbf{X}^{train} \mathbf{W} - \mathbf{Y}^{train} \right\|_2^2 = 0$$

$\Rightarrow$

$$\mathbf{W} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Normal  
equation

Only works  
analytically for linear  
regression

# Analytical learning in Linear Regression

**Performance measure P** : mean square error (**MSE**)  $\text{MSE}_{test}$  :

$$\text{MSE}_{test} = \frac{1}{N_{test}} \sum_{n=1}^{N_{test}} (\hat{y}_i^{test} - y_i^{test})^2 = \frac{1}{N_{test}} \left\| \hat{\mathbf{Y}}^{test} - \mathbf{Y}^{test} \right\|_2^2 = \frac{1}{N_{test}} \left\| \mathbf{X}^{test} \mathbf{W} - \mathbf{Y}^{test} \right\|_2^2$$

Parameters  $\mathbf{W} \in \mathbb{R}^D$  are found by minimizing  $\text{MSE}_{train}$

$$\text{MSE}_{train} = \frac{1}{N_{train}} \sum_{n=1}^{N_{train}} (\hat{y}_i^{train} - y_i^{train})^2 = \frac{1}{N_{train}} \left\| \hat{\mathbf{Y}}^{train} - \mathbf{Y}^{train} \right\|_2^2 = \frac{1}{N_{train}} \left\| \mathbf{X}^{train} \mathbf{W} - \mathbf{Y}^{train} \right\|_2^2$$

For optimal weights  $\mathbf{W}$ , the gradient of  $\text{MSE}_{train}$  should be 0:

$$\nabla_{\mathbf{W}} \text{MSE}_{train} = \nabla_{\mathbf{W}} \frac{1}{N_{train}} \left\| \hat{\mathbf{Y}}^{train} - \mathbf{Y}^{train} \right\|_2^2 = \frac{1}{N_{train}} \nabla_{\mathbf{W}} \left\| \mathbf{X}^{train} \mathbf{W} - \mathbf{Y}^{train} \right\|_2^2 = 0$$

$\Rightarrow$

$$\mathbf{W} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Normal equation

Only works  
analytically for linear  
regression

Complexity  $O(N^3)$  to  
 $O(N^{2.4})$

# Gradient Descent for Linear Regression

**Performance measure P :** mean square error (**MSE**)  $\text{MSE}_{test}$  :

$$\text{MSE}_{test} = \frac{1}{N_{test}} \sum_{n=1}^{N_{test}} (\hat{y}_i^{test} - y_i^{test})^2 = \frac{1}{N_{test}} \left\| \hat{\mathbf{Y}}^{test} - \mathbf{Y}^{test} \right\|_2^2 = \frac{1}{N_{test}} \left\| \mathbf{X}^{test} \mathbf{W} - \mathbf{Y}^{test} \right\|_2^2$$

Parameters  $\mathbf{W} \in \mathbb{R}^D$  are found by minimizing  $\text{MSE}_{train}$

$$\text{MSE}_{train} = \frac{1}{N_{train}} \sum_{n=1}^{N_{train}} (\hat{y}_i^{train} - y_i^{train})^2 = \frac{1}{N_{train}} \left\| \hat{\mathbf{Y}}^{train} - \mathbf{Y}^{train} \right\|_2^2 = \frac{1}{N_{train}} \left\| \mathbf{X}^{train} \mathbf{W} - \mathbf{Y}^{train} \right\|_2^2$$

For optimal weights  $\mathbf{W}$ , the gradient of  $\text{MSE}_{train}$  should be 0:

$$\nabla_{\mathbf{W}} \text{MSE}_{train} = \nabla_{\mathbf{W}} \frac{1}{N_{train}} \left\| \hat{\mathbf{Y}}^{train} - \mathbf{Y}^{train} \right\|_2^2 = \frac{1}{N_{train}} \nabla_{\mathbf{W}} \left\| \mathbf{X}^{train} \mathbf{W} - \mathbf{Y}^{train} \right\|_2^2 = 0$$


$$\Rightarrow \mathbf{W} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Only works  
analytically for linear  
regression

Complexity  $O(N^3)$  to  
 $O(N^{2.4})$

# Gradient Descent


Optimize weights **W** **gradually** by moving against the gradients of  $\text{MSE}_{\text{train}}$





# Gradient Descent

Optimize weights **W** **gradually** by moving against the gradients of  $\text{MSE}_{\text{train}}$




$$\theta_1 = \theta_0 - \eta g'_\theta(\theta_0)$$

$$\theta_2 = \theta_1 - \eta g'_\theta(\theta_1)$$

$\vdots$

# Gradient Descent


Optimize weights **W** **gradually** by moving against the gradients of  $\text{MSE}_{\text{train}}$



$$\theta_1 = \theta_0 - \eta g'_\theta(\theta_0)$$


$$\theta_2 = \theta_1 - \eta g'_\theta(\theta_1)$$

$\vdots$




# Choice of learning rate

Optimize weights **W** **gradually** by moving against the gradients of  $\text{MSE}_{\text{train}}$



# Choice of learning rate

Optimize weights **W** **gradually** by moving against the gradients of  $\text{MSE}_{\text{train}}$




Learning rate

$$\theta_1 = \theta_0 - \eta g'_\theta(\theta_0)$$


$$\theta_2 = \theta_1 - \eta g'_\theta(\theta_1)$$

$\vdots$



# Choice of learning rate

Optimize weights **W** **gradually** by moving against the gradients of  $\text{MSE}_{\text{train}}$




Learning rate

$$\theta_1 = \theta_0 - \eta g'_\theta(\theta_0)$$

$$\theta_2 = \theta_1 - \eta g'_\theta(\theta_1)$$

$\vdots$




Local minimum

Global minimum

# Choice of learning rate

Optimize weights **W** **gradually** by moving against the gradients of  $\text{MSE}_{\text{train}}$




Learning rate

$$\theta_1 = \theta_0 - \eta g'_\theta(\theta_0)$$

$$\theta_2 = \theta_1 - \eta g'_\theta(\theta_1)$$

$\vdots$



# Control question

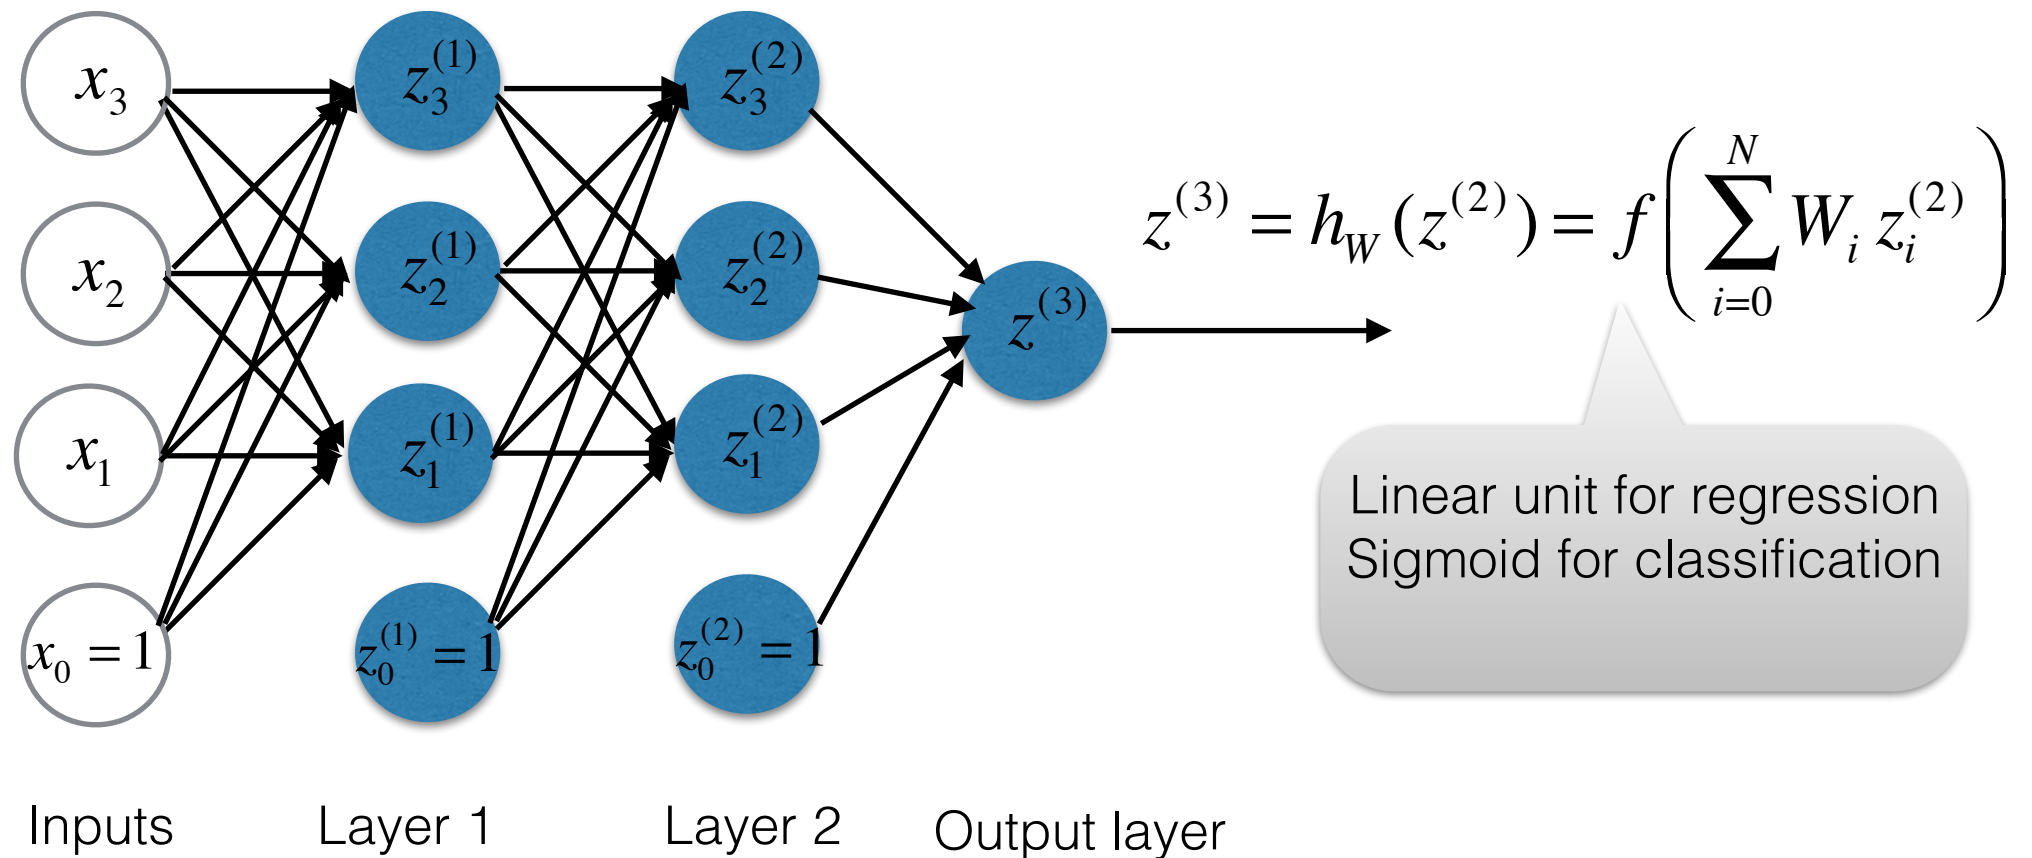
Q: Select all correct statements:

1. Gradient Descent always leads to a unique solution starting from any initial point, no matter what the objective function is.
2. Gradient Descent has one free parameter called the learning rate.
3. A good choice of the learning rate is important: if the learning rate is too small, it takes long for the algorithm to converge, but if it is too high, the algorithm may diverge.
4. Making the learning rate variable (larger initially in the training, and smaller as the training progresses) may accelerate ML algorithms.
5. The fastest way to find the optimal learning rate for a ML algorithm is to simply add it as one more model parameter and optimize over it, along with other model parameters, in the process of minimization of the train error.

Correct answers: 2, 3, 4.

# Gradient Descent for Neural Networks

- Gradient descent can be applied not only to one neuron, but to a neural network too:



- Training by backpropagation = reverse-mode autodiff